

MediaPipe and DINOv2 Representations for Hand-Strength Prediction in Broadcast Poker

Minu Choi
School of Data Science
University of Virginia
Charlottesville, VA
qce2dp@virginia.edu

Abstract—Televised poker commentary and a recent broadcast “AI bluff detector” popularized the claim that nonverbal behavior reveals hand strength on camera. We test that claim with a fully automated pipeline over 46.6 hours of publicly broadcast cash-game footage with exposed hole cards: eight sessions, 12,417 decisions, 9,801 of them with equity labels. Game state is read from the broadcast graphics by OCR, hands are reassembled into decision windows, and hand strength is labeled by Monte Carlo equity. MediaPipe face and body streams are summarized into behavioral features for camera-covered decisions of tracked players. Our central methodological finding is an identity-contamination artifact. Under progressively stricter face-attribution hygiene, an apparent behavioral effect first strengthened to a nominally resolvable Δ AUC of +0.096 (95% CI [+0.005, +0.181]), then vanished entirely (+0.009, [−0.058, +0.074]) once every enrolled reference face was verified. The spurious effect was produced by frames in which a lookalike neighbor was silently credited to the tracked player. A preregistered replication on two additional sessions, with the analysis frozen before download, decisively failed: pooled across four sessions, face and pose features change out-of-session AUC by −0.075 (95% CI [−0.134, −0.020], $n = 292$ covered decisions), with approximately 91% power to detect the registered +0.096. Sensitivity analyses show the degradation is not estimator variance. No regularization scheme rescues the features; a behavior-only model is anti-predictive out of session (AUC 0.349, below chance under a hand-level permutation test, $p = 0.018$); the effect persists within a single player across his four sessions. A second preregistered round resolved the within-session question on all eight sessions: behavior-only AUC is 0.587 (95% CI [0.515, 0.655], $n = 494$), excluding chance under the frozen criterion, while cross-session transfer stays at or below chance (0.430, [0.358, 0.502]). A preregistered frozen DINOv2 probe adds nothing over betting action out of session (+0.012, [−0.079, +0.103]), extending the conclusion to learned representations. Tells exist and are machine-readable within a session, and they do not transfer across sessions. Decision timing, which needs no camera, is bounded below +0.01 AUC over betting action. We conclude that identity hygiene, not feature engineering, is the binding methodological risk for behavioral video datasets built on multi-person footage, and that behavioral hand-strength models of this class do not generalize across sessions.

Index Terms—nonverbal behavior, poker, deception detection, person re-identification, preregistration, broadcast video analysis, negative results

I. INTRODUCTION

Poker folklore holds that players leak the strength of their hands through nonverbal “tells,” and the 2026 World Series of Poker broadcast popularized the idea that such leakage

is machine-readable, overlaying a behavioral hand-strength model on live coverage. The scientific evidence behind that idea is thinner than its television presence suggests. The best controlled result reports that naive observers judge professional players’ hand quality above chance from arm movements alone, at modest effect sizes, while judgments from the face are at or below chance [1]. Automated deception detection benchmarks report within-domain accuracies in the sixties and cross-domain transfer barely above chance [2]. No prior work, to our knowledge, has tested the broadcast-poker claim end to end: does adding automatically extracted nonverbal behavior improve prediction of hand strength over the betting action alone, evaluated out of session, on public footage?

This paper contributes four things.

(1) **A reproducible pipeline** that converts raw broadcast video into a labeled decision-level dataset. The stages are OCR of the broadcast graphics package into per-second game-state snapshots, a state machine that reassembles hands and decision windows, Monte Carlo equity labels from the exposed hole cards, and per-decision behavioral features from face landmarks, blendshapes, and body pose. Features carry per-player normalization and explicit bet-size leakage controls. A released demo renderer overlays every extracted signal on the source broadcast (Figure 1).

(2) **An identity-contamination case study** (Section VI-A). Behavioral datasets built on multi-person footage must attribute every measured face to the right person. We document three small engineering gaps: an open-set similarity threshold, nondeterministic reference enrollment, and unbounded nearest-position matching. Together they let a lookalike neighbor’s face be silently measured as the tracked player’s. The resulting artifact produced an apparently resolvable behavioral effect that grew under partial cleaning and vanished under full cleaning (Figure 2). We distill the fix into a verification protocol.

(3) **A preregistered, adequately powered replication failure** (Section VI-B). With the analysis frozen before the new footage was downloaded, face and pose features reduce pooled out-of-session AUC by 0.075 (95% CI [−0.134, −0.020]). The study had approximately 91% power to detect the registered positive estimate.

(4) **A session-locality finding** (Section VI-C). Behavioral features are not merely uninformative out of session: a

behavior-only model scores AUC 0.349, significantly below chance under a hand-level permutation test ($p = 0.018$), and the inversion persists within the same player across his four sessions. Within sessions the same features carry real signal, resolved in a second preregistered round on eight sessions (pooled 0.587, CI [0.515, 0.655], excluding chance). Learned behavior-to-strength associations are session-specific in sign: real enough to fit, wrong enough to mislead on a new night. Tells are session-local, at both ends of the claim.

II. RELATED WORK

Poker AI plays cards, not people. Superhuman poker agents reached and then surpassed top human play from game state and betting action alone [3], [4]. No behavioral channel appears anywhere in that line of work. Whether human nonverbal behavior carries residual, machine-readable signal on top of the betting action is exactly the claim the broadcast overlay popularized, and the one we test.

Behavioral cues to poker hand strength. Slepian et al. [1] showed observers rate hand quality above chance from short clips of arm motion (smoothness of the chip push), while face-only clips were directionally below chance. This motivates our smoothness features and our expectation that any true effect is small. Work on automated deception detection in the wild reports within-domain accuracy around 60 to 67 percent on gameshow and courtroom corpora and severe degradation across domains [2]. Cross-session failure of this kind is dataset shift [5]: the deployment distribution differs from training in ways the model cannot see. Our result extends this pattern to a naturalistic, high-stakes, repeated-game setting with objective ground truth, and sharpens it to shift across nights within one venue, one broadcast package, and one player.

Person identification hygiene. Face verification [6] and person re-identification [7] under compression, occlusion, and pose variation are well studied; sunglasses are endemic at poker tables. Deciding whether a detected face belongs to an enrolled identity at all is the classic open-set problem [8], and our initial fixed-threshold gate is exactly the naive open-set design that literature warns against. What is underappreciated is the role of attribution error as a silent noise channel in downstream behavioral datasets. Label errors are known to destabilize benchmark conclusions [9]; our artifact is the feature-side analogue, attribution noise manufacturing a positive finding out of another person’s face.

Methodological safeguards. We use preregistration [10], grouped bootstrap confidence intervals [11], and leakage review of engineered features. Trajectory smoothness metrics follow [12], [13]; keypoint trajectories are One-Euro filtered [14] before any derivative is taken; face and pose streams come from MediaPipe [15]; OCR uses PP-OCR [16].

III. DATA

Footage. Eight complete Hustler Casino Live cash-game sessions (October 2025 to July 2026, 46.6 hours total) featuring one high-volume professional (all eight sessions) and a second recreational player (one session). Round 1 used four of

them: the February, April, and both July (a) and (b) broadcasts of 2026. The October, December, March, and July (c) sessions were added under the round 2 preregistration. The stream exposes hole cards in its graphics, providing ground truth without any inference from behavior.

IV. METHODOLOGY

A. Preprocessing and Game-State Extraction

Region-cropped OCR reads the pot, stakes, board, and per-player panels (name, stack, action text, amount to call) once per second. Card sprites are template-matched with a saturation-gated segmentation robust to the two background styles in our footage. A state machine folds snapshots into hands and decision windows: a window opens when a player’s panel shows an amount to call and closes at their action flip. Cutaway gaps are re-merged when positions and hole cards agree, and a preroll filter drops inset replays of older sessions. Table I summarizes the result: 1,461 hands and 12,417 decisions across the eight sessions, 9,801 of them with equity labels. Labels are Monte Carlo equity versus a random hand with deterministic seeds, binned into six strength classes; `is_weak` marks decisions taken with equity below 0.45, and `is_bluff` marks aggressive such decisions.

B. Behavioral Feature Extraction

For each decision window of a mapped player, frames are searched for the player’s face (Section VI-A). Accepted frames contribute blendshape series: blink rate, smile asymmetry and mean, gaze dispersion, and brow activity. A pose crop that follows the accepted face contributes posture and wrist kinematics: lean variability, head motion, normalized wrist peak speed, spectral arc length, and negative log dimensionless jerk, which is amplitude-invariant by construction and therefore cannot encode bet size. Six event detectors (downward-gaze rate, hand near face, freeze fraction and longest freeze, chip-shuffle periodicity, forward lean) are rates, fractions, ratios, or shoulder-width geometry, and passed the same leakage review. Earlier two-session analyses showed the detectors carry no signal, so the preregistration excluded them from the primary test; Section VI-C confirms the null on the pooled data. All features are z-scored per player. Coverage is honest and low: the primary player’s face passes identity gating in 10 to 33 percent of window frames depending on session. In the round 1 sessions, 292 labeled decisions meet the 0.5 coverage threshold; across all eight sessions, 494 do.

C. Evaluation Protocol

Leave-one-session-out (LOSO): models train on all other sessions and predict the held-out session; predictions are pooled. Baselines use betting features only (position, sizing ratios, pot odds, street, prior raises, stack-to-pot, timing). Deltas are reported with hand-grouped bootstrap 95% CIs (2,000 resamples), resampling hands so correlated decisions within a hand never narrow the interval.



Fig. 1. One annotated frame from the released demo renderer (October 2025 session, hand 120, river decision). Green box: identity-gated face attribution with its re-identification score. Colored contours are MediaPipe face landmarks (brows, eyelids, lips); the magenta ray is the gaze vector, the yellow arrow the head direction, and the yellow trace the One-Euro-filtered acting-wrist path that feeds the kinematic features; the tag beside the box is a live event detector (brow furrowed). Top right: the held-out six-class strength distribution, the is_bluff probability with and without behavioral features, and live per-window readouts standardized against the player’s own baseline. Bottom center: hand identifier, street, pot, actor, and decision clock from the OCR game state. The gold player panels and pot graphics are part of the source broadcast. This session is the one in which the primary player wore no eyewear, so eyelid and gaze tracking are unoccluded.

TABLE I

DATASET SUMMARY. A DECISION IS COVERED WHEN AT LEAST HALF OF ITS WINDOW FRAMES PASS IDENTITY GATING (FACE) OR POSE TRACKING FOR A MAPPED PLAYER; ONLY COVERED, LABELED DECISIONS ENTER THE ABLATIONS. ROUND 1 COMPRISES FEB, APR, JUL (A), AND JUL (B) 2026; THE OTHER FOUR SESSIONS WERE ADDED IN ROUND 2.

Session	Hours	Snapshots	Hands	Decisions	Labeled	Covered
Oct 2025	6.7	22,636	156	1,736	1,079	75
Dec 2025	6.4	22,774	196	1,814	1,223	26
Feb 2026	5.2	17,501	141	1,320	1,163	114
Mar 2026	6.4	21,851	212	1,770	1,511	31
Apr 2026	5.5	18,814	164	1,188	1,027	55
Jul 2026 (a)	5.2	18,436	225	1,561	1,278	51
Jul 2026 (b)	5.9	20,730	192	1,570	1,347	72
Jul 2026 (c)	5.3	18,011	175	1,458	1,173	70
Total	46.6	160,753	1,461	12,417	9,801	494

D. Identity Attribution

Attributing faces on a multi-person broadcast requires deciding, per frame, whether a detected face is the tracked player. Our initial design used an open-set rule: accept the first face whose equalized grayscale patch correlates at least 0.55 with any enrolled reference patch of the player, a threshold chosen from a measured same-player/cross-player separation of 0.74

versus 0.25. Section VI-A documents how three gaps in this design manufactured a spurious finding; the corrected protocol is: (i) enroll references deterministically, one detector instance per reference frame; (ii) bound enrollment by position and refuse specifications whose intended face is not found; (iii) verify every enrolled patch by eye; (iv) sweep each session for faces that score above threshold against the references, and register every non-target such face as a distractor; (v) attribute a face only if it clears the threshold and beats every distractor score by a margin, best candidate wins; (vi) treat staff and commentators as candidate distractors too (one of ours scored 0.60).

V. MODELS

Every model in the pipeline was chosen under three constraints. The corpus is large in hours but small in supervised examples: hundreds of camera-covered decisions, not tens of thousands. The footage is compressed 1080p broadcast video, so perception components must tolerate compression, cutaways, and occlusion. And the study’s claims rest on auditability, so each stage must produce outputs a human can verify.

Game-state reading. Broadcast graphics are read by the PP-OCRv6 detection and recognition models (tiny tier) of

the PP-OCR family [16], applied to fixed regions so the recognizer sees one panel at a time. The graphics package is a fixed template, so the problem is throughput rather than layout understanding: per-second sampling of 46.6 hours yields 160,753 snapshots, all processed on CPU. Card sprites are identified by normalized template correlation instead of a learned detector because the sprite set is closed and rendered deterministically, which makes template matching both exact and inspectable.

Behavioral streams. Face geometry, blendshapes, and body keypoints come from the MediaPipe Face Landmarker and Pose Landmarker (lite) models [15]. We chose MediaPipe because its blendshape outputs are named, interpretable facial signals (blink, smile, brow, gaze) that map directly onto the tells discussed in the poker literature, and because both landmarkers run at broadcast frame rates on CPU, which the 46.6-hour corpus requires.

Identity matching. Faces are attributed by correlation of equalized grayscale patches against enrolled references rather than by a deep face embedding such as ArcFace [6]. The choice is deliberate: every acceptance is a comparison against a reference patch a human has verified, and the artifact study of Section VI-A depends on being able to inspect exactly what the matcher saw. A stronger embedding would widen the match margin but faces the same open-set decisions that Section IV-D addresses by protocol.

Hand-strength labels. Equity is computed by Monte Carlo simulation against a random opponent hand using the `treys` evaluator with deterministic seeds. Equity versus a random hand ignores opponents’ actual ranges; we accept that simplification because it requires no modeling assumptions and is computed identically across sessions, so it cannot introduce session-specific label bias.

Analysis models. All primary tests use logistic regression (`scikit-learn` [17], $C = 1.0$) on standardized, median-imputed features, and ridge regression for the continuous-equity arm. A linear probe is the standard instrument for asking whether a representation carries signal, and with $n = 292$ to 494 covered decisions, anything more flexible would mostly measure its own variance. Stronger and sparser regularization variants (L2 at $C = 0.1$ and 0.01, L1 at $C = 0.1$ and 0.05) are reported as sensitivity arms in Section VI-C.

Learned representation. Round 2 tests whether hand-crafted summaries are the bottleneck by replacing them with DINOv2 ViT-S/14 embeddings [18], kept frozen. Identity-gated face crops and their pose crops are embedded per frame, pooled to the per-window mean and standard deviation, and reduced by PCA to 32 components fit on training folds only. DINOv2 was chosen as a strong general-purpose self-supervised representation that transfers without fine-tuning; freezing it keeps the probe protocol identical to the hand-crafted arm and avoids fitting tens of millions of parameters to hundreds of examples.




Fig. 2. An apparent behavioral effect tracks attribution pollution, not behavior. Points are Δ AUC of face and pose features over the betting baseline for `is_weak` with hand-grouped bootstrap 95% CIs. The first three estimates are in-sample over two sessions under progressively stricter identity hygiene; the fourth is the preregistered four-session test.

VI. RESULTS

A. The Identity-Contamination Artifact

Three gaps in the initial identity design of Section IV-D turned face attribution into a contamination channel.

Gap 1: lookalikes above threshold. A neighboring player correlated 0.66 with the tracked player’s references. Whenever the tracked player’s face dipped out of detectability (chin on hand, looking down), the neighbor was accepted in his place. In one session, 82 of 193 windows contained some wrong-person frames; one window was entirely the neighbor.

Gap 2: nondeterministic enrollment. Reference patches were enrolled by seeking to hand-verified timestamps and taking the detected face nearest a specified position. Detection under a warm video-mode tracker is stateful: which faces are found at a seek depends on the seeks before it. Re-runs enrolled different faces from the same specification.

Gap 3: unbounded nearest-position matching. When the intended player went undetected at a reference timestamp, the nearest detected face was enrolled with no distance bound. In two of four sessions this placed a neighbor’s face inside the player’s positive reference library, which also neutralized the distractor defense added after Gap 1: the enrolled lookalike outscored his own distractor reference.

Figure 2 shows the consequence. With no identity checks the apparent effect was $+0.044$ $[-0.028, +0.117]$. After registering the known lookalike as a distractor, but with a wrong face still enrolled as a positive reference, the effect strengthened to $+0.096$ $[+0.005, +0.181]$, nominally resolvable. With enrollment made deterministic (fresh tracker per reference frame), position-bounded (a face must lie within 110 px of its specification), and every enrolled patch visually verified, the in-sample effect collapsed to $+0.009$ $[-0.058, +0.074]$. The mechanism is mundane: another person’s face, measured during the tracked player’s decisions, covaries with decision context (cameras cut to neighbors during long tanks) in ways that imitate a tell.




Fig. 3. No analysis choice rescues the features (round 1: pooled LOSO over the four round 1 sessions). All arms are Δ AUC over the betting baseline for `is_weak` with hand-grouped bootstrap 95% CIs.

B. Preregistered Replication

Before downloading any new footage we froze, in a committed preregistration document: the hypothesis (face and pose features improve pooled LOSO AUC for `is_weak`), the registered estimate (+0.096), the exact pipeline settings, the seat-mapping discipline for new sessions (commit-moment identification and a distractor audit before any model output), the pass criterion (pooled CI excluding zero on the positive side), and a prohibition on post-hoc adjustment. Two sessions were selected by roster and layout compatibility; one candidate was excluded before extraction because it used a different graphics package, and the exclusion was logged in the pre-registration.

The pooled four-session test returned

$$\Delta\text{AUC} = -0.075, \quad 95\% \text{ CI } [-0.134, -0.020], \quad n = 292.$$

The CI excludes zero on the negative side: added behavioral features reduce out-of-session performance. Under a normal approximation to the bootstrap distribution ($\text{SE} \approx 0.029$), the test had approximately 91% power to detect the registered +0.096, and 80% power for any delta of at least +0.082. The upper confidence bound rules out all positive effects at the 2.5% level. The betting-only baseline itself scores AUC 0.549 pooled out of session, against 0.75 within session, replicating the known brutality of cross-session transfer.

C. Model Results: Why Negative, Not Just Null

A significantly negative delta invites the objection that the estimator, not the signal, is at fault: ten added dimensions can hurt a small model even when they carry information. Four analyses reject this reading (Figure 3).

Regularization does not help. Stronger L2 shrinkage leaves the delta unchanged or worse ($C = 0.01$: -0.115), and sparse L1 selection is worst of all (-0.154): selecting the few behavioral features that look best in training amplifies reliance on associations that do not hold in the held-out session.

Behavior alone is anti-predictive. A behavior-only model scores AUC 0.349, far below chance. Whatever association the model learns between behavior and weakness in three sessions,




Fig. 4. The below-chance transfer is significant (round 1). Hand-level label permutations within each session, re-running the full cross-session pipeline per draw (500 draws), place the observed behavior-only AUC in the null’s lower tail ($p = 0.018$).

the reverse tends to hold in the fourth. Under label exchange this would be signal; under session exchange it is instability.

The inversion survives within a person. Restricted to the primary player’s four sessions (one identity, one seat style, $n = 280$), the delta is -0.069 [$-0.128, -0.009$]. Idiosyncrasy across people does not explain the failure; instability within a person across nights does.

Session geometry explains only part. Re-standardizing features per player-session (removing session-level offsets induced by camera geometry, since a lean angle or apparent motion scale depends on the shot) moves the behavior-only AUC from 0.349 to 0.387 and the delta to -0.056 [$-0.115, -0.004$]. Session offsets are therefore a real confound channel that future work must normalize away, and still not the whole story.

The inversion is significant; within-session signal was marginal in round 1. Two further analyses locate the failure precisely. First, a permutation test that shuffles labels at the hand level within each session and re-runs the full cross-session pipeline places the observed behavior-only AUC of 0.349 (hand-grouped bootstrap CI [0.279, 0.429], entirely below chance) in the lower tail of its null (null mean 0.476, SD 0.058, $p = 0.018$; Figure 4); transfer is not merely absent but systematically inverted. Second, within-session grouped cross-validation (5-fold by hand, per session, pooled over the four round 1 sessions) gives a behavior-only AUC of 0.574 with a hand-grouped bootstrap CI of [0.492, 0.652] and per-session values from 0.418 to 0.676; Section VI-E resolves this estimate on all eight sessions (Figure 5). Together these say the model is not hallucinating. It finds real within-session structure in at least some sessions, but the sign of that structure is session-specific (Table II: no feature keeps its fitted sign across all four round 1 sessions), so exporting it is worse than ignoring behavior. This is the paper’s clearest positive characterization: behavioral hand-strength associations in this setting are session-local.

The event detectors are a true null, not an inver-




Fig. 5. Session locality, resolved on eight sessions (round 2). Within-session grouped cross-validation finds real but heterogeneous behavioral signal (pooled CI excludes chance); the same features transfer at or below chance across sessions.

ision. As secondary arms on the pooled four-session data, the six leakage-proofed event detectors alone give $+0.011$ $[-0.037, +0.063]$, and the union of all behavioral features gives -0.060 $[-0.128, +0.005]$. The contrast with face and pose is instructive: the low-dimensional detectors, built as rates and ratios with explicit leakage review, neither help nor hurt, while the richer face and pose summaries are the source of the inversion.

D. Decision Timing and Continuous Targets

Two further analyses use the label information more efficiently. **Decision timing at full power.** Time to act needs no camera, so it is available for all 4,815 labeled decisions of every player in the round 1 sessions, not just the 292 camera-covered ones. Added to the betting baseline it changes *is_weak* AUC by $+0.004$ (95% CI $[-0.002, +0.010]$) and *is_bluff* by -0.000 $[-0.005, +0.004]$: the most cited poker tell, tank duration, is bounded below one hundredth of AUC in this setting. **Continuous equity.** Regressing Monte Carlo equity directly (ridge, out-of-session Spearman correlation) uses the full label rather than a threshold. The face and pose degradation replicates and tightens ($\Delta\rho = -0.099$, CI $[-0.144, -0.047]$ on covered rows), while timing shows the study’s only resolvable positive effect, at negligible magnitude ($\Delta\rho = +0.005$, CI $[+0.002, +0.008]$, $n = 4,815$). The latter matters mainly as a sensitivity check: the pipeline and evaluation resolve genuinely small effects when the sample allows, so the behavioral nulls are not an artifact of blunt instruments.

TABLE II
NO BEHAVIORAL FEATURE KEEPS ITS COEFFICIENT SIGN ACROSS ALL FOUR ROUND 1 SESSIONS (BEHAVIOR-ONLY LOGISTIC MODEL PER SESSION, *IS_WEAK*, STANDARDIZED FEATURES; $+/-$ IS THE FITTED SIGN).

Feature	Feb	Apr	Jul (a)	Jul (b)
blink rate	+	−	+	+
smile asymmetry	−	−	−	+
smile mean	+	−	+	−
gaze dispersion	+	−	+	+
brow activity	+	+	+	−
wrist peak speed	+	+	−	−
wrist smoothness (LDJ)	+	+	−	+
wrist SPARC	+	+	−	−
lean variability	−	−	−	+
head motion	−	−	+	−

E. Round 2: Preregistered Resolution and Learned Representations

A second preregistration, committed before any round 2 footage was downloaded, fixed two experiments: four additional sessions to resolve the within-session estimate, and the frozen DINOv2 probe of Section V to test whether a learned representation changes the out-of-session answer.

Within-session signal is real. Pooled within-session grouped cross-validation across all eight sessions gives a behavior-only AUC of 0.587 with a hand-grouped bootstrap 95% CI of $[0.515, 0.655]$ ($n = 494$), excluding chance under the frozen resolution criterion (Figure 5). Per-session values remain strongly heterogeneous (0.26 to 0.76). The best session (0.756) is the one in which the primary player wore no sunglasses, though the second bare-faced session scores low at small n , so the occlusion account stays a hypothesis. The cross-session arm on the same enlarged pool stays at or below chance (behavior-only 0.430, CI $[0.358, 0.502]$; face and pose delta over betting -0.034 , CI $[-0.080, +0.012]$). The session-locality claim is thereby resolved at both ends: tells exist within a night and do not transfer across nights.

A learned representation does not change the answer. The preregistered DINOv2 probe adds $+0.012$ over the betting baseline out of session (CI $[-0.079, +0.103]$, $n = 257$ embedding-covered decisions): a null, extending the conclusion from hand-crafted features to a standard frozen representation. Two nuances are worth recording. The embedding probe alone transfers at 0.545, above the inverted hand-crafted features, so embeddings are uninformative rather than misleading out of session. Within sessions the probe scores 0.398 (CI $[0.320, 0.476]$), below chance where the ten hand-crafted features find real signal. At these sample sizes, a 1536-dimensional representation reduced to 32 components is less stable than ten purpose-built features, a capacity lesson for small behavioral corpora.

VII. DISCUSSION

Three readings of these results deserve separation. The narrow reading is engineering: with tabular summary features over MediaPipe streams, a linear model, eight sessions, and camera-limited coverage, out-of-session behavioral signal is

not extractable, and the attempt costs accuracy. The middle reading is statistical. Effects of the size the literature supports, a few hundredths of AUC, require an order of magnitude more covered decisions than 46.6 hours of broadcast yields. Any study in this regime that reports a large positive behavioral effect should therefore be audited first for identity contamination, which we show can manufacture exactly such an effect. The broad reading is substantive. Even within one professional, behavior-strength associations invert across nights, which is what one expects if experienced players' surface behavior is dominated by context, table dynamics, and deliberate randomization rather than stable leakage. Round 2 resolves the distinction that round 1 could not: tells exist (within-session signal excludes chance under a preregistered criterion) and are contextual and unstable across nights, which entails that broadcast-scale models of this class cannot use them out of session.

The artifact deserves the emphasis we give it. It produced, in our own hands, a preregistration-worthy positive estimate with a confidence interval excluding zero, out of nothing but face misattribution. It grew stronger under partial cleaning because the partially cleaned pipeline retained the most correlated contamination. Behavioral-video findings built on multi-person footage without an explicit attribution audit should be treated as provisional.

VIII. LIMITATIONS

One primary player carries most covered decisions; sunglasses make eye features noise for him, as for many poker subjects. Coverage is bound by broadcast camera direction, not by us. Card-reading errors survive at a low rate (rank confusions on small sprites) and add label noise to equity. The betting baseline is deliberately simple; a stronger baseline would only shrink headroom for behavior. Our features are window-level summaries. Sequence models over raw streams remain untested here, though they face the same coverage and identity constraints with more capacity to overfit them. Two round 2 observations mark future work rather than conclusions. The best within-session score came from the one session in which the primary player wore no sunglasses, so occlusion may bound the face channel; a targeted comparison of bare-faced versus occluded sessions would test that directly. And the DINOv2 probe scored below chance within sessions where ten purpose-built features found signal, suggesting that representation capacity at this sample size, not representation quality, is what binds; the remedy is more covered decisions, not bigger models.

IX. CONCLUSION

We built an end-to-end pipeline from broadcast poker video to labeled decision-level behavioral data and asked whether nonverbal behavior adds predictive power over betting action for hand strength. Under a preregistered and adequately powered test, the answer is that it subtracts out of session. Behavior-strength associations are learnable within a session, but their signs are session-specific, so models that carry

them to a new session perform significantly below chance. Along the way we showed how easily the opposite conclusion is manufactured: a single misattributed face in a reference library produced a nominally significant positive effect that survived partial cleaning. The attribution-hygiene protocol, the preregistrations, and the full analysis history are released with the code. For behavioral video research, the order of operations this study recommends is: verify identity first, preregister second, and only then believe an effect.

X. ETHICS AND REPRODUCIBILITY

All footage is publicly broadcast with player consent to the stream's hole-card exposure. We analyze it post hoc and redistribute no footage; the single annotated frame in Figure 1 is reproduced from the public broadcast to illustrate the system. We publish code, derived features, timestamps, and the preregistrations rather than clips. Pipeline, analysis code, seat maps, audit tooling, figure scripts, and the frozen preregistrations with their logged results are available at <https://github.com/minuuva/Poker-Bluff-Detector> (private during review, public with publication). The full analysis history, including the artifact's discovery and every estimate in Figure 2, is preserved in that repository's commit log.

REFERENCES

- [1] M. L. Slepian, S. G. Young, A. M. Rutchick, and N. Ambady, "Quality of professional poker players' poker hands is perceivable from their movements," *Psychological Science*, vol. 24, no. 11, pp. 2335–2338, 2013.
- [2] X. Guo, N. M. Selvaraj, Z. Yu, A. W.-K. Kong, B. Shen, and A. Kot, "Audio-visual deception detection: DOLOS dataset and parameter-efficient crossmodal learning," in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2023.
- [3] N. Brown and T. Sandholm, "Superhuman AI for heads-up no-limit poker: Libratus beats top professionals," *Science*, vol. 359, no. 6374, pp. 418–424, 2018.
- [4] —, "Superhuman AI for multiplayer poker," *Science*, vol. 365, no. 6456, pp. 885–890, 2019.
- [5] J. Quiñero-Candela, M. Sugiyama, A. Schwaighofer, and N. D. Lawrence, Eds., *Dataset Shift in Machine Learning*. MIT Press, 2009.
- [6] J. Deng, J. Guo, N. Xue, and S. Zafeiriou, "Arcface: Additive angular margin loss for deep face recognition," in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019.
- [7] L. Zheng, Y. Yang, and A. G. Hauptmann, "Person re-identification: Past, present and future," *arXiv preprint arXiv:1610.02984*, 2016.
- [8] W. J. Scheirer, A. Rocha, A. Sapkota, and T. E. Boult, "Toward open set recognition," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 35, no. 7, pp. 1757–1772, 2013.
- [9] C. G. Northcutt, A. Athalye, and J. Mueller, "Pervasive label errors in test sets destabilize machine learning benchmarks," in *Proceedings of the NeurIPS Track on Datasets and Benchmarks*, 2021.
- [10] B. A. Nosek, C. R. Ebersole, A. C. DeHaven, and D. T. Mellor, "The preregistration revolution," *Proceedings of the National Academy of Sciences*, vol. 115, no. 11, pp. 2600–2606, 2018.
- [11] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1993.
- [12] S. Balasubramanian, A. Melendez-Calderon, A. Roby-Brami, and E. Burdet, "On the analysis of movement smoothness," *Journal of NeuroEngineering and Rehabilitation*, vol. 12, no. 1, p. 112, 2015.
- [13] N. Hogan and D. Sternad, "Sensitivity of smoothness measures to movement duration, amplitude, and arrests," *Journal of Motor Behavior*, vol. 41, no. 6, pp. 529–534, 2009.
- [14] G. Casiez, N. Roussel, and D. Vogel, "1 euro filter: A simple speed-based low-pass filter for noisy input in interactive systems," in *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems (CHI)*, 2012, pp. 2527–2530.

- [15] C. Lugaresi, J. Tang, H. Nash, C. McClanahan, E. Uboweja, M. Hays, F. Zhang, C.-L. Chang, M. G. Yong, J. Lee, W.-T. Chang, W. Hua, M. Georg, and M. Grundmann, “Mediapipe: A framework for building perception pipelines,” *arXiv preprint arXiv:1906.08172*, 2019.
- [16] Y. Du, C. Li, R. Guo, X. Yin, W. Liu, J. Zhou, Y. Bai, Z. Yu, Y. Yang, Q. Dang, and H. Wang, “Pp-ocr: A practical ultra lightweight OCR system,” *arXiv preprint arXiv:2009.09941*, 2020.
- [17] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [18] M. Oquab, T. Darcet, T. Moutakanni, H. V. Vo, M. Szafraniec, V. Khaidov, P. Fernandez, D. Haziza, F. Massa, A. El-Nouby, M. Assran, N. Ballas, W. Galuba, R. Howes, P.-Y. Huang, S.-W. Li, I. Misra, M. Rabbat, V. Sharma, G. Synnaeve, H. Xu, H. Jégou, J. Mairal, P. Labatut, A. Joulin, and P. Bojanowski, “DINOv2: Learning robust visual features without supervision,” *Transactions on Machine Learning Research*, 2024.